

Update on Ara

01/06/2022

Matteo Perotti

Matheus Cavalcante

Nils Wistoff

Gianmarco Ottavi

Professor Luca Benini


Integrated Systems Laboratory

ETH Zürich

FFT timeline

- ✓ Python golden model
- ✓ Scalar DIT + DIF (CVA6 only)
- Vectorized DIF algorithm
 - ✓ First complete implementation
 - Debugging
- ✗ Performance analysis
- ✗ Optimization
- ✗ ISA extension?

Vectorized FFT - DIF



Vectorized FFT - Issues

- Problem with intrinsics
 - Typed
 - Missing *mask-vector slides*

Vectorized FFT - Issues

- Problem with intrinsics
 - Typed
 - Missing *mask-vector slides*



- Load mask vectors from memory
- Check if bottleneck

Vectorized FFT - Issues

- Ara hangs
 - First complex algorithm
 - All the units are stressed together

Vectorized FFT - Issues

- Ara hangs
 - First complex algorithm
 - All the units are stressed together



- Debugging

Vector FFT - DIF

```
// Butterfly until the end
for (unsigned int i = 1; i < log2_nfft; ++i) {
    // Bump the twiddle pointers.
    twiddles_re += vl;
    twiddles_im += vl;

    // Load twiddle factors
    twiddle_re = vle32_v_f32m1(twiddles_re, vl);
    twiddle_im = vle32_v_f32m1(twiddles_im, vl);

    // HALVE vl_mask for permutation stage
    vl_mask >>= 1;

    // Create the current mask level
    //vslideup_vx_f32m1(mask_vec_buf, mask_vec, 0, vl_mask);
    //mask_vec = vmxor_mm_b32(mask_vec, mask_vec_buf, vl);
    mask_vec = vlm_v_b32(mask_addr_vec[i], vl);
    mask_vec_buf = vmnot_m_b32(mask_vec, vl);

    // 1) Get the upper wing output
    vbuf_re = vfadd_vv_f32m1(upper_wing_re, lower_wing_re, vl);
    vbuf_im = vfadd_vv_f32m1(upper_wing_im, lower_wing_im, vl);
    // 2) Get the lower wing output
    lower_wing_re = vfsub_vv_f32m1(upper_wing_re, lower_wing_re, vl);
    lower_wing_im = vfsub_vv_f32m1(upper_wing_im, lower_wing_im, vl);
    // Copy labels
    upper_wing_re = vbuf_re;
    upper_wing_im = vbuf_im;
    // 3) Multiply lower wing for the twiddle factor
    vbuf_re = cmplx_mul_re_vv(lower_wing_re, lower_wing_im, twiddle_re, twiddle_im, vl);
    lower_wing_im = cmplx_mul_im_vv(lower_wing_re, lower_wing_im, twiddle_re, twiddle_im, vl);
    lower_wing_re = vbuf_re; // Just for the label. Verify that there is no actual copy of this vector

    // Different permutation for the last round
    if (i != log2_nfft - 1) {
        // Permutate the numbers
        vbuf_re = vsliedown_vx_f32m1_m(mask_vec_buf, vbuf_re, upper_wing_re, vl/2, vl/2);
        vbuf_im = vsliedown_vx_f32m1_m(mask_vec_buf, vbuf_im, upper_wing_im, vl/2, vl/2);
        upper_wing_re = vsliedup_vx_f32m1(upper_wing_re, lower_wing_re, vl/2, vl/2);
        upper_wing_im = vsliedup_vx_f32m1(upper_wing_im, lower_wing_im, vl/2, vl/2);
        lower_wing_re = vmerge_vvm_f32m1(mask_vec, vbuf_re, lower_wing_re, vl/2);
        lower_wing_im = vmerge_vvm_f32m1(mask_vec, vbuf_im, lower_wing_im, vl/2);
    }
}
```


Vectorized FFT - DIF

```
// Butterfly until the end
for (unsigned int i = 1; i < log2_nfft; ++i) {
    // Bump the twiddle pointers.
    twiddles_re += vl;
    twiddles_im += vl;

    // Load twiddle factors
    twiddle_re = vle32_v_f32m1(twiddles_re, vl);
    twiddle_im = vle32_v_f32m1(twiddles_im, vl);

    // HALVE vl_mask for permutation stage
    vl_mask >>= 1;

    // Create the current mask level
    //vslideup_vx_f32m1(mask_vec_buf, mask_vec, 0, vl_mask);
    //mask_vec = vmxor_mm_b32(mask_vec, mask_vec_buf, vl);
    mask_vec = vlm_v_b32(mask_addr_vec[i], vl);
    mask_vec_buf = vmnot_m_b32(mask_vec, vl);

    // 1) Get the upper wing output
    vbuf_re = vfadd_vv_f32m1(upper_wing_re, lower_wing_re, vl);
    vbuf_im = vfadd_vv_f32m1(upper_wing_im, lower_wing_im, vl);
    // 2) Get the lower wing output
    lower_wing_re = vfsub_vv_f32m1(upper_wing_re, lower_wing_re, vl);
    lower_wing_im = vfsub_vv_f32m1(upper_wing_im, lower_wing_im, vl);
    // Copy labels
    upper_wing_re = vbuf_re;
    upper_wing_im = vbuf_im;
    // 3) Multiply lower wing for the twiddle factor
    vbuf_re = cmplx_mul_re_vv(lower_wing_re, lower_wing_im, twiddle_re, twiddle_im, vl);
    lower_wing_im = cmplx_mul_im_vv(lower_wing_re, lower_wing_im, twiddle_re, twiddle_im, vl);
    lower_wing_re = vbuf_re; // Just for the label. Verify that there is no actual copy of this vector

    // Different permutation for the last round
    if (i != log2_nfft - 1) {
        // Permutate the numbers
        vbuf_re = vsldown_vx_f32m1_m(mask_vec_buf, vbuf_re, upper_wing_re, vl/2, vl/2);
        vbuf_im = vsldown_vx_f32m1_m(mask_vec_buf, vbuf_im, upper_wing_im, vl/2, vl/2);
        upper_wing_re = vsldown_vx_f32m1(upper_wing_re, lower_wing_re, vl/2, vl/2);
        upper_wing_im = vsldown_vx_f32m1(upper_wing_im, lower_wing_im, vl/2, vl/2);
        lower_wing_re = vmerge_vvm_f32m1(mask_vec, vbuf_re, lower_wing_re, vl/2);
        lower_wing_im = vmerge_vvm_f32m1(mask_vec, vbuf_im, lower_wing_im, vl/2);
    }
}
```

Load the twiddle factors from memory
For this first implementation, they are already in memory
Reuse can be improved in next implementations

Vectorized FFT - DIF

```
// Butterfly until the end
for (unsigned int i = 1; i < log2_nfft; ++i) {
    // Bump the twiddle pointers.
    twiddles_re += vl;
    twiddles_im += vl;

    // Load twiddle factors
    twiddle_re = vle32_v_f32m1(twiddles_re, vl);
    twiddle_im = vle32_v_f32m1(twiddles_im, vl);

    // HALVE vl_mask for permutation stage
    vl_mask >>= 1;

    // Create the current mask level
    //vslideup_vx_f32m1(mask_vec_buf, mask_vec, 0, vl_mask);
    //mask_vec = vmxor_mm b32(mask_vec, mask_vec_buf, vl);
    mask_vec = vlm_v_b32(mask_addr_vec[i], vl);
    mask_vec_buf = vmnot_m_b32(mask_vec, vl);

    // 1) Get the upper wing output
    vbuf_re = vfadd_vv_f32m1(upper_wing_re, lower_wing_re, vl);
    vbuf_im = vfadd_vv_f32m1(upper_wing_im, lower_wing_im, vl);
    // 2) Get the lower wing output
    lower_wing_re = vfsub_vv_f32m1(upper_wing_re, lower_wing_re, vl);
    lower_wing_im = vfsub_vv_f32m1(upper_wing_im, lower_wing_im, vl);
    // Copy labels
    upper_wing_re = vbuf_re;
    upper_wing_im = vbuf_im;
    // 3) Multiply lower wing for the twiddle factor
    vbuf_re = cmplx_mul_re_vv(lower_wing_re, lower_wing_im, twiddle_re, twiddle_im, vl);
    lower_wing_im = cmplx_mul_im_vv(lower_wing_re, lower_wing_im, twiddle_re, twiddle_im, vl);
    lower_wing_re = vbuf_re; // Just for the label. Verify that there is no actual copy of this vector

    // Different permutation for the last round
    if (i != log2_nfft - 1) {
        // Permutate the numbers
        vbuf_re = vslidedown_vx_f32m1_m(mask_vec_buf, vbuf_re, upper_wing_re, vl/2, vl/2);
        vbuf_im = vslidedown_vx_f32m1_m(mask_vec_buf, vbuf_im, upper_wing_im, vl/2, vl/2);
        upper_wing_re = vslideup_vx_f32m1(upper_wing_re, lower_wing_re, vl/2, vl/2);
        upper_wing_im = vslideup_vx_f32m1(upper_wing_im, lower_wing_im, vl/2, vl/2);
        lower_wing_re = vmerge_vvm_f32m1(mask_vec, vbuf_re, lower_wing_re, vl/2);
        lower_wing_im = vmerge_vvm_f32m1(mask_vec, vbuf_im, lower_wing_im, vl/2);
    }
}
```

**Load mask vector for permutation from memory
(or, commented: create the mask vector within Ara)**

Vectorized FFT - DIF

```
// Butterfly until the end
for (unsigned int i = 1; i < log2_nfft; ++i) {
    // Bump the twiddle pointers.
    twiddles_re += vl;
    twiddles_im += vl;

    // Load twiddle factors
    twiddle_re = vle32_v_f32m1(twiddles_re, vl);
    twiddle_im = vle32_v_f32m1(twiddles_im, vl);

    // HALVE vl_mask for permutation stage
    vl_mask >>= 1;

    // Create the current mask level
    //vslideup_vx_f32m1(mask_vec_buf, mask_vec, 0, vl_mask);
    //mask_vec = vmxor_mm_b32(mask_vec, mask_vec_buf, vl);
    mask_vec = vlm_v_b32(mask_addr_vec[i], vl);
    mask_vec_buf = vmnnot_m_b32(mask_vec, vl);

    // 1) Get the upper wing output
    vbuf_re = vfadd_vv_f32m1(upper_wing_re, lower_wing_re, vl);
    vbuf_im = vfadd_vv_f32m1(upper_wing_im, lower_wing_im, vl);
    // 2) Get the lower wing output
    lower_wing_re = vfsub_vv_f32m1(upper_wing_re, lower_wing_re, vl);
    lower_wing_im = vfsub_vv_f32m1(upper_wing_im, lower_wing_im, vl);
    // Copy labels
    upper_wing_re = vbuf_re;
    upper_wing_im = vbuf_im;
    // 3) Multiply lower wing for the twiddle factor
    vbuf_re = cmplx_mul_re_vv(lower_wing_re, lower_wing_im, twiddle_re, twiddle_im, vl);
    lower_wing_im = cmplx_mul_im_vv(lower_wing_re, lower_wing_im, twiddle_re, twiddle_im, vl);
    lower_wing_re = vbuf_re; // Just for the label. Verify that there is no actual copy of this vector

    // Different permutation for the last round
    if (i != log2_nfft - 1) {
        // Permutate the numbers
        vbuf_re = vsldown_vx_f32m1_m(mask_vec_buf, vbuf_re, upper_wing_re, vl/2, vl/2);
        vbuf_im = vsldown_vx_f32m1_m(mask_vec_buf, vbuf_im, upper_wing_im, vl/2, vl/2);
        upper_wing_re = vslideup_vx_f32m1(upper_wing_re, lower_wing_re, vl/2, vl/2);
        upper_wing_im = vslideup_vx_f32m1(upper_wing_im, lower_wing_im, vl/2, vl/2);
        lower_wing_re = vmerge_vvm_f32m1(mask_vec, vbuf_re, lower_wing_re, vl/2);
        lower_wing_im = vmerge_vvm_f32m1(mask_vec, vbuf_im, lower_wing_im, vl/2);
    }
}
```

Butterfly the vectors

Vectorized FFT - DIF

```
// Butterfly until the end
for (unsigned int i = 1; i < log2_nfft; ++i) {
    // Bump the twiddle pointers.
    twiddles_re += vl;
    twiddles_im += vl;

    // Load twiddle factors
    twiddle_re = vle32_v_f32m1(twiddles_re, vl);
    twiddle_im = vle32_v_f32m1(twiddles_im, vl);

    // HALVE vl_mask for permutation stage
    vl_mask >>= 1;

    // Create the current mask level
    //vslideup_vx_f32m1(mask_vec_buf, mask_vec, 0, vl_mask);
    //mask_vec = vmxor_mm_b32(mask_vec, mask_vec_buf, vl);
    mask_vec = vlm_v_b32(mask_addr_vec[i], vl);
    mask_vec_buf = vmnot_m_b32(mask_vec, vl);

    // 1) Get the upper wing output
    vbuf_re = vfadd_vv_f32m1(upper_wing_re, lower_wing_re, vl);
    vbuf_im = vfadd_vv_f32m1(upper_wing_im, lower_wing_im, vl);
    // 2) Get the lower wing output
    lower_wing_re = vfsub_vv_f32m1(upper_wing_re, lower_wing_re, vl);
    lower_wing_im = vfsub_vv_f32m1(upper_wing_im, lower_wing_im, vl);
    // Copy labels
    upper_wing_re = vbuf_re;
    upper_wing_im = vbuf_im;
    // 3) Multiply lower wing for the twiddle factor
    vbuf_re = cmplx_mul_re_vv(lower_wing_re, lower_wing_im, twiddle_re, twiddle_im, vl);
    lower_wing_im = cmplx_mul_im_vv(lower_wing_re, lower_wing_im, twiddle_re, twiddle_im, vl);
    lower_wing_re = vbuf_re; // Just for the label. Verify that there is no actual copy of this vector

    // Different permutation for the last round
    if (i != log2_nfft - 1) {
        // Permutate the numbers
        vbuf_re = vslidedown_vx_f32m1_m(mask_vec_buf, vbuf_re, upper_wing_re, vl/2, vl/2);
        vbuf_im = vslidedown_vx_f32m1_m(mask_vec_buf, vbuf_im, upper_wing_im, vl/2, vl/2);
        upper_wing_re = vslideup_vx_f32m1(upper_wing_re, lower_wing_re, vl/2, vl/2);
        upper_wing_im = vslideup_vx_f32m1(upper_wing_im, lower_wing_im, vl/2, vl/2);
        lower_wing_re = vmerge_vvm_f32m1(mask_vec, vbuf_re, lower_wing_re, vl/2);
        lower_wing_im = vmerge_vvm_f32m1(mask_vec, vbuf_im, lower_wing_im, vl/2);
    }
}
```

Permute the vectors for the “next step”
 $4 * (vl/2)$ -long slides -> costly

FFT timeline

- ✓ Python golden model
- ✓ Scalar DIT + DIF (CVA6 only)
- Vectorized DIF algorithm
 - ✓ First complete implementation
 - Debugging
- ✗ Performance analysis
- ✗ Optimization
- ✗ ISA extension?